

CS241 -- Lecture Notes: B-Trees

[Daisy Tang](#)

[Back To Lectures Notes](#)

This lecture covers Section 10.2 of our text book and [here](#).

Introduction to B-Trees

A B-tree is a tree data structure that keeps data sorted and allows searches, insertions, and deletions in logarithmic amortized time. Unlike self-balancing binary search trees, it is optimized for systems that read and write large blocks of data. It is most commonly used in database and file systems.


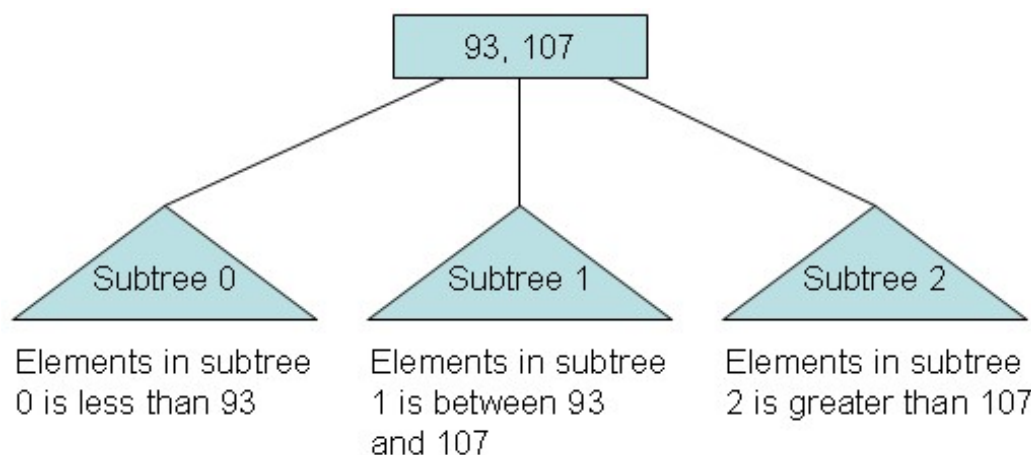
The B-Tree Rules

Important properties of a B-tree:

- B-tree nodes have many more than two children.
- A B-tree node may contain more than just a single element.

The set formulation of the B-tree rules: Every B-tree depends on a positive constant integer called **MINIMUM**, which is used to determine how many elements are held in a single node.

- **Rule 1:** The root can have as few as one element (or even no elements if it also has no children); every other node has at least **MINIMUM** elements.
- **Rule 2:** The maximum number of elements in a node is twice the value of **MINIMUM**.
- **Rule 3:** The elements of each B-tree node are stored in a partially filled array, sorted from the smallest element (at index 0) to the largest element (at the final used position of the array).
- **Rule 4:** The number of subtrees below a nonleaf node is always one more than the number of elements in the node.
 - Subtree 0, subtree 1, ...
- **Rule 5:** For any nonleaf node:
 1. An element at index i is greater than all the elements in subtree number i of the node, and
 2. An element at index i is less than all the elements in subtree number $i + 1$ of the node.
- **Rule 6:** Every leaf in a B-tree has the same depth. Thus it ensures that a B-tree avoids the problem of an unbalanced tree.



The Set Class Implementation with B-Trees

Remember that "Every child of a node is also the root of a smaller B-tree".

```
public class IntBalancedSet implements Cloneable
{
    private static final int MINIMUM = 200;
    private static final int MAXIMUM = 2*MINIMUM;
    int dataCount;
    int[] data = new int[MAXIMUM + 1];
    int childCount;
    IntBalancedSet[] subset = new IntBalancedSet[MAXIMUM + 2];

    // Constructor: initialize an empty set
    public IntBalancedSet()

    // add: add a new element to this set, if the element was already in the set, then
    // there is no change.
    public void add(int element)

    // clone: generate a copy of this set.
    public IntBalancedSet clone()

    // contains: determine whether a particular element is in this set
    public boolean contains(int target)


    // remove: remove a specified element from this set
    public boolean remove(int target)
}
```

Searching for a Target in a Set

The psuedocode:

1. Make a local variable, i , equal to the first index such that $\text{data}[i] \geq \text{target}$. If there is no such index, then set i equal to dataCount , indicating that none of the elements is greater than or equal to the target.
2. if (we found the target at $\text{data}[i]$)
 return true;
 else if (the root has no children)
 return false;
 else return $\text{subset}[i].\text{contains}(\text{target})$;

See the following example, try to search for 10.




We can implement a private method:

- private int firstGE(int target), which returns the first location in the root such that $\text{data}[x] \geq \text{target}$. If there's no such location, then return value is dataCount .

Adding an Element to a B-Tree

It is easier to add a new element to a B-tree if we relax one of the B-tree rules.


Loose addition allows the root node of the B-tree to have $\text{MAXIMUM} + 1$ elements. For example, suppose we want to add 18 to the tree:



The above result is an illegal B-tree. Our plan is to perform a loose addition first, and then fix the root's problem.

The Loose Addition Operation for a B-Tree:


```
private void looseAdd(int element)
{
    1. i = firstGE(element) // find the first index such that data[i] >= element
    2. if (we found the new element at data[i]) return; // since there's already a copy
    in the set
    3. else if (the root has no children)
        Add the new element to the root at data[i]. (shift array)
    4. else {
        subset[i].looseAdd(element);
        if the root of subset[i] now has an excess element, then fix that problem
        before returning.
    }
}
```




```
private void fixExcess(int i)
// precondition: (i < childCount) and the entire B-tree is valid except that subset[i]
// has MAXIMUM + 1 elements.
// postcondition: the tree is rearranged to satisfy the loose addition rule
```

Fixing a Child with an Excess Element:


- To fix a child with $\text{MAXIMUM} + 1$ elements, the child node is split into two nodes that each contain MINIMUM elements. This leaves one extra element, which is passed up to the parent.
- It is always the *middle* element of the split node that moves upward.
- The parent of the split node gains one additional child and one additional element.
- The children of the split node have been equally distributed between the two smaller nodes.





Fixing the Root with an Excess Element:

- Create a new root.
- `fixExcess(0)`.



Removing an Element from a B-Tree

Loose removal rule: *Loose removal allows to leave a root that has one element too few.*

```
public boolean remove(int target)
{
    answer = looseRemove(target);
    if ((dataCount == 0) && (childCount == 1))
        Fix the root of the entire tree so that it no longer has zero elements;
    return answer;
}

private boolean looseRemove(int target)
{
    1. i = firstGE(target)
    2. Deal with one of these four possibilities:
        2a. if (root has no children and target not found) return false.
        2b. if (root has no children but target found) {
            remove the target
            return true
        }
        2c. if (root has children and target not found) {
            answer = subset[i].looseRemove(target)
            if (subset[i].dataCount < MINIMUM)
                fixShortage(i)
            return true
        }
        2d. if (root has children and target found) {
            data[i] = subset[i].removeBiggest()
            if (subset[i].dataCount < MINIMUM)
                fixShortage(i)
            return true
        }
    }
}


private void fixShortage(int i)
// Precondition: (i < childCount) and the entire B-tree is valid except that subset[i]
has MINIMUM - 1 elements.
```

```
// Postcondition: problem fixed based on the looseRemoval rule.
```

```
private int removeBiggest()
```

```
// Precondition: (dataCount > 0) and this entire B-tree is valid
```

```
// Postcondition: the largest element in this set has been removed and returned. The  
entire B-tree is still valid based on the looseRemoval rule.
```




Fixing Shortage in a Child:

When $\text{fixShortage}(i)$ is activated, we know that $\text{subset}[i]$ has $\text{MINIMUM} - 1$ elements. There are four cases that we need to consider:

Case 1: Transfer an extra element from $\text{subset}[i-1]$. Suppose $\text{subset}[i-1]$ has more than the MINIMUM number of elements.


- Transfer $\text{data}[i-1]$ down to the front of $\text{subset}[i].\text{data}$.
- Transfer the final element of $\text{subset}[i-1].\text{data}$ up to replace $\text{data}[i-1]$.
- If $\text{subset}[i-1]$ has children, transfer the final child of $\text{subset}[i-1]$ over to the front of $\text{subset}[i]$.




Case 2: Transfer an extra element from $\text{subset}[i+1]$. Suppose $\text{subset}[i+1]$ has more than the MINIMUM number of elements.

Case 3: Combine $\text{subset}[i]$ with $\text{subset}[i-1]$. Suppose $\text{subset}[i-1]$ has only MINIMUM elements.

- Transfer $\text{data}[i-1]$ down to the end of $\text{subset}[i-1].\text{data}$.
- Transfer all the elements and children from $\text{subset}[i]$ to the end of $\text{subset}[i-1]$.
- Disconnect the node $\text{subset}[i]$ from the B-tree by shifting $\text{subset}[i+1]$, $\text{subset}[i+2]$ and so on leftward.





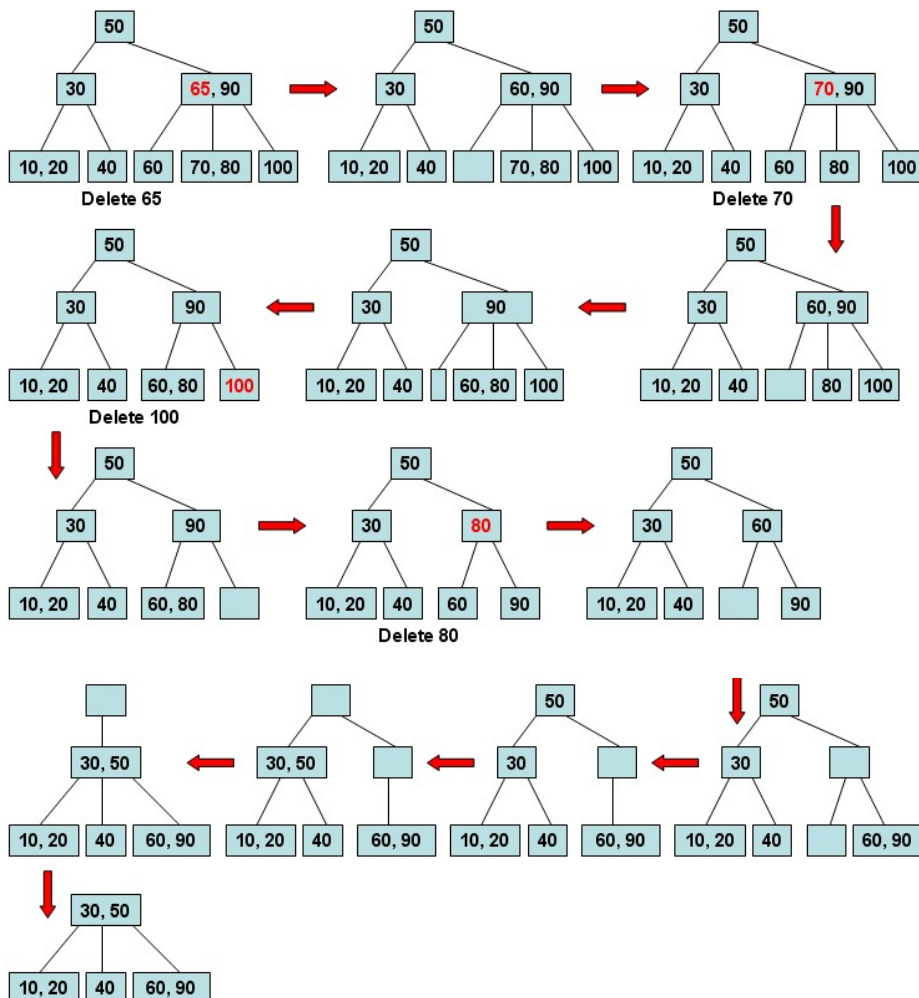
Case 4: Combine $\text{subset}[i]$ with $\text{subset}[i+1]$. Suppose $\text{subset}[i+1]$ has only MINIMUM elements.

We may need to continue activating `fixShortage()` until the B-tree rules are satisfied.

Removing the Biggest Element from a B-Tree:


```
private int removeBiggest()
{
    if (root has no children)
        remove and return the last element
    else {
        answer = subset[childCount-1].removeBiggest()
        if (subset[childCount-1].dataCount < MINIMUM)
            fixShortage(childCount-1)
        return answer
    }
}
```

A more concrete example for node deletion:



2-3 Tree ([applet](#))

A 2-3 tree is a type of B-tree where every node with children (internal node) has either two children and one data element (2-nodes) or three children and two data elements (3-node). Leaf nodes have no children and one or two data elements.



Practice: insert the following values in sequence to a 2-3 tree: 50, 19, 21, 66, 84, 29, and 54.

Practice: now, delete 50, 66, 54 from the above B-tree.

Take a look at this [2-3 tree animation](#).

Exercises:

Build a 2-3 tree with the following ten values: 10, 9, 8, ..., and 1. Delete 10, 8 and 6 from the tree.

Trees -- Time Analysis

The implementation of a B-tree is efficient since the depth of the tree is kept small.

Worst-case times for tree operations: the worst-case time performance for the following operations are all $O(d)$, where d is the depth of the tree:

1. Adding an element to a binary search tree (BST), a heap, or a B-tree.
2. Removing an element from a BST, a heap, or a B-tree.
3. Searching for a specified element in a BST or a B-tree.

Time Analysis for BST

Suppose a BST has n elements. What is the maximum depth the tree could have?

- A BST with n elements could have a depth as big as $n-1$.

Worst-Case Times for BSTs:

- Adding an element, removing an element, or searching for an element in a BST with n elements is $O(n)$.

Time Analysis for Heaps

Remember that a heap is a complete BST, so each level must be full before proceeding to the next level.

Number of nodes needed for a heap to reach depth d is: $(1 + 2 + 4 + 8 + \dots + 2^{d-1}) + 1 = 2^d = n$. Thus $d = \log_2 n$.

Worst-Case Times for Heap Operations:

- Adding or removing an element in a heap with n elements is $O(\log n)$.

Time Analysis for B-Tree

Suppose a B-tree has n elements and M is the maximum number of children a node can have. What is the maximum depth the tree could have? What is the minimum depth the tree could have?

- The worst-case depth (maximum depth) of a B-tree is: $\log_{M/2} n$.
- The best-case depth (minimum depth) of a B-tree is: $\log_M n$.

Worst-Case Times for B-Trees:

- Adding or removing an element in a B-tree with n elements is $O(\log n)$.

Learning Objectives

When you complete this section, you will be able to:

- list the rules for a B-tree and determine whether a tree satisfies these rules.
- do a simulation by hand of the algorithms for searching, inserting, and removing an element from a B-tree.
- use the B-tree data structure to implement a set class.
- use Java's DefaultMutableTreeNode and JTree classes in simple programs that use trees.
- understand the complexities of each operation.

Last updated: Oct. 2012