

# CS350: Data Structures

## B-Trees

---

James Moscola

Department of Engineering & Computer Science

York College of Pennsylvania



# Introduction

---

- **All of the data structures that we've looked at thus far have been memory-based data structures**
- **What if we have more data than we can store in main system memory?**
- **What if we need to store a data structure on disk?**
  - Are the data structures we've talked about thus far suitable for disk?
  - What features of previous data structures might need to be altered?

# Memory Access vs. Disk Access

---

- **Following pointers in memory is a relatively quick operation when compared to following pointers on disk**
- **Hard disk drives are very slow**
  - A 7200 RPM disk drive rotates 7200 times in one minute (each revolution takes  $1/120$  of a second ... ~8.3 ms)
  - On average, must spin the disk half way to get to the desired data
  - For this example, can do about 120 disk accesses per second
  - A 3 GHz processor can compute the result of ~3 Billion instructions per second (6 Billion if dual core)
  - A disk access is about 25 Million times SLOWER than a modern processor (i.e. the processor can execute 25 Million instructions in the time that it takes one disk access to finish)

# Memory Access vs. Disk Access


---

- **Storing large amounts of data in a binary tree on disk could take very long to access**
  - Given a perfectly balanced binary tree, the number of disk accesses to search for an element on disk is  $O(\log N)$ 
    - May take on the order of several seconds
  - Given an unbalanced binary tree (worst case) the number of disk accesses to search for an element on disk is  $O(N)$ 
    - **Example: A tree with 10 Million nodes stored on disk**
      - May require 10 Million disk accesses to find last node. At 8.3ms/disk access, that's 83 million ms (~23 hours)

# B-Tree Idea

---

- **Want to reduce the number of disk accesses to a small number (i.e. 3 or 4)**
  - Increase the fanout of a tree (i.e. increase the # of children at each node)
- **Reduce the height of the tree**
  - Reduces the number of disk accesses required to traverse tree
- **Requires extra computation, but processor time is very cheap compared to disk access time**



# B-Trees

---

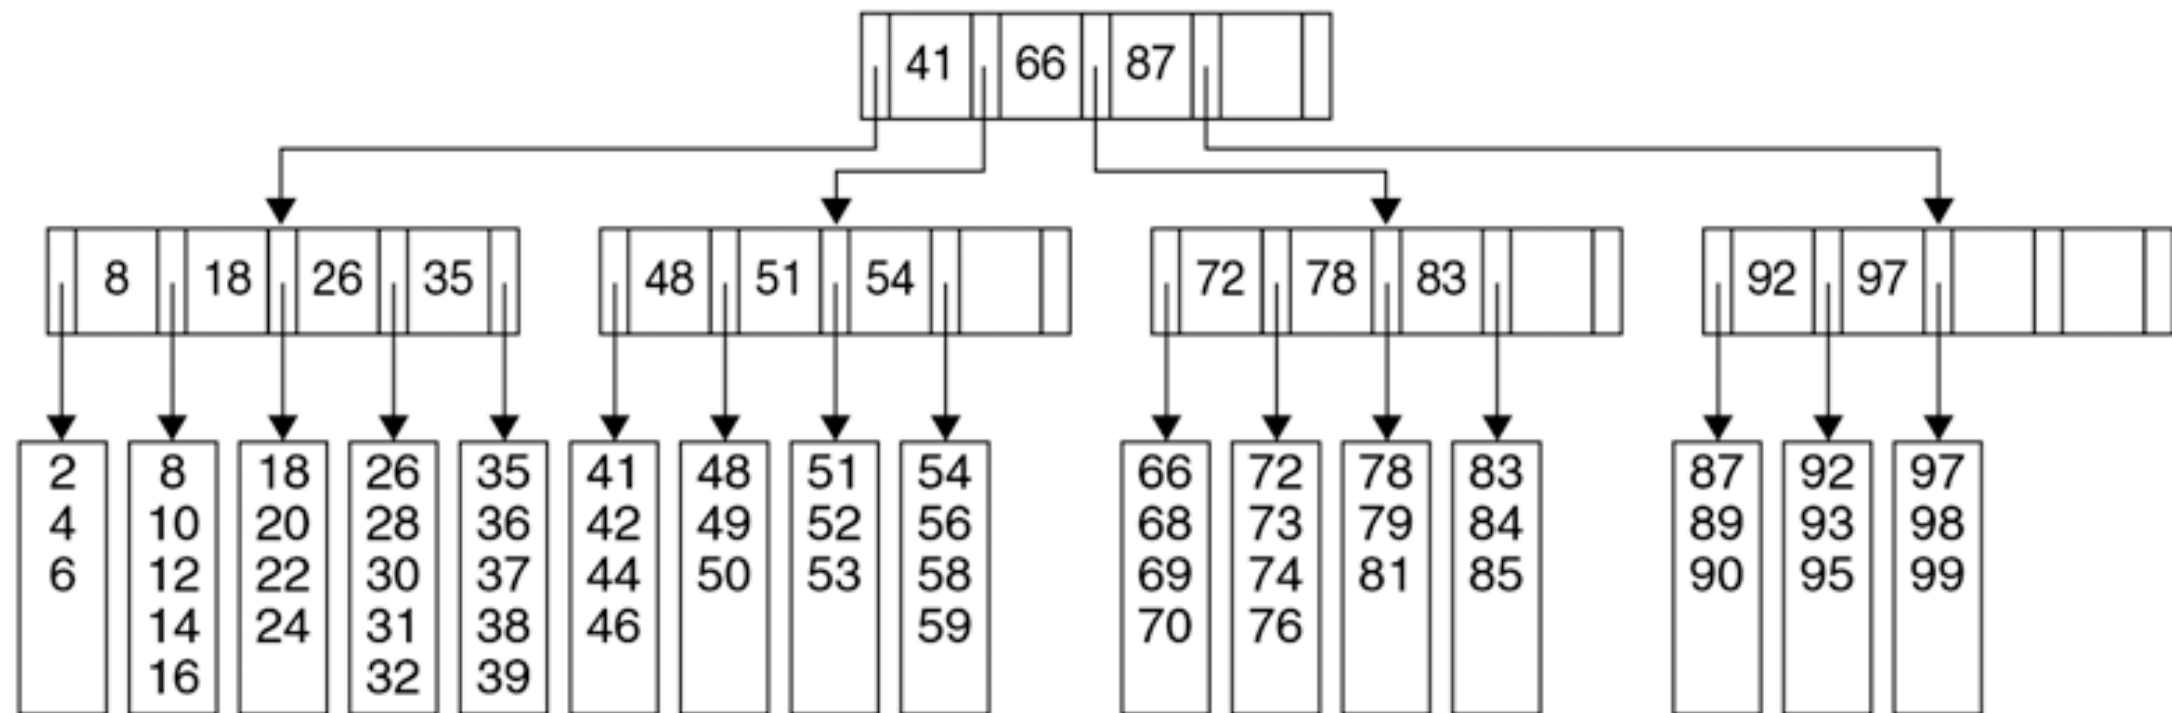
- **The number of children that a node may have is  $M$  (an  $M$ -ary tree)**
- **Each node requires  $M-1$  keys to determine which branch to take when traversing the tree**
  - In a binary tree  $M=2$ , and therefore it requires only 1 key at each node to determine which child pointer to follow

# B-Tree Properties

---

- **A B-tree of order  $M$  is an  $M$ -ary tree with the following properties:**
  - (1) Data items are stored at leaf nodes
  - (2) Non-leaf nodes store as many as  $M-1$  keys to guide the searching process; key  $i$  represents the smallest key in subtree  $i+1$
  - (3) The root is either a leaf or has between 2 and  $M$  children
  - (4) All non-leaf nodes (except the root) have between  $\lceil M/2 \rceil$  and  $M$  children
  - (5) All leaf nodes are at the same depth and have between  $\lceil L/2 \rceil$  and  $L$  data items, for some value  $L$

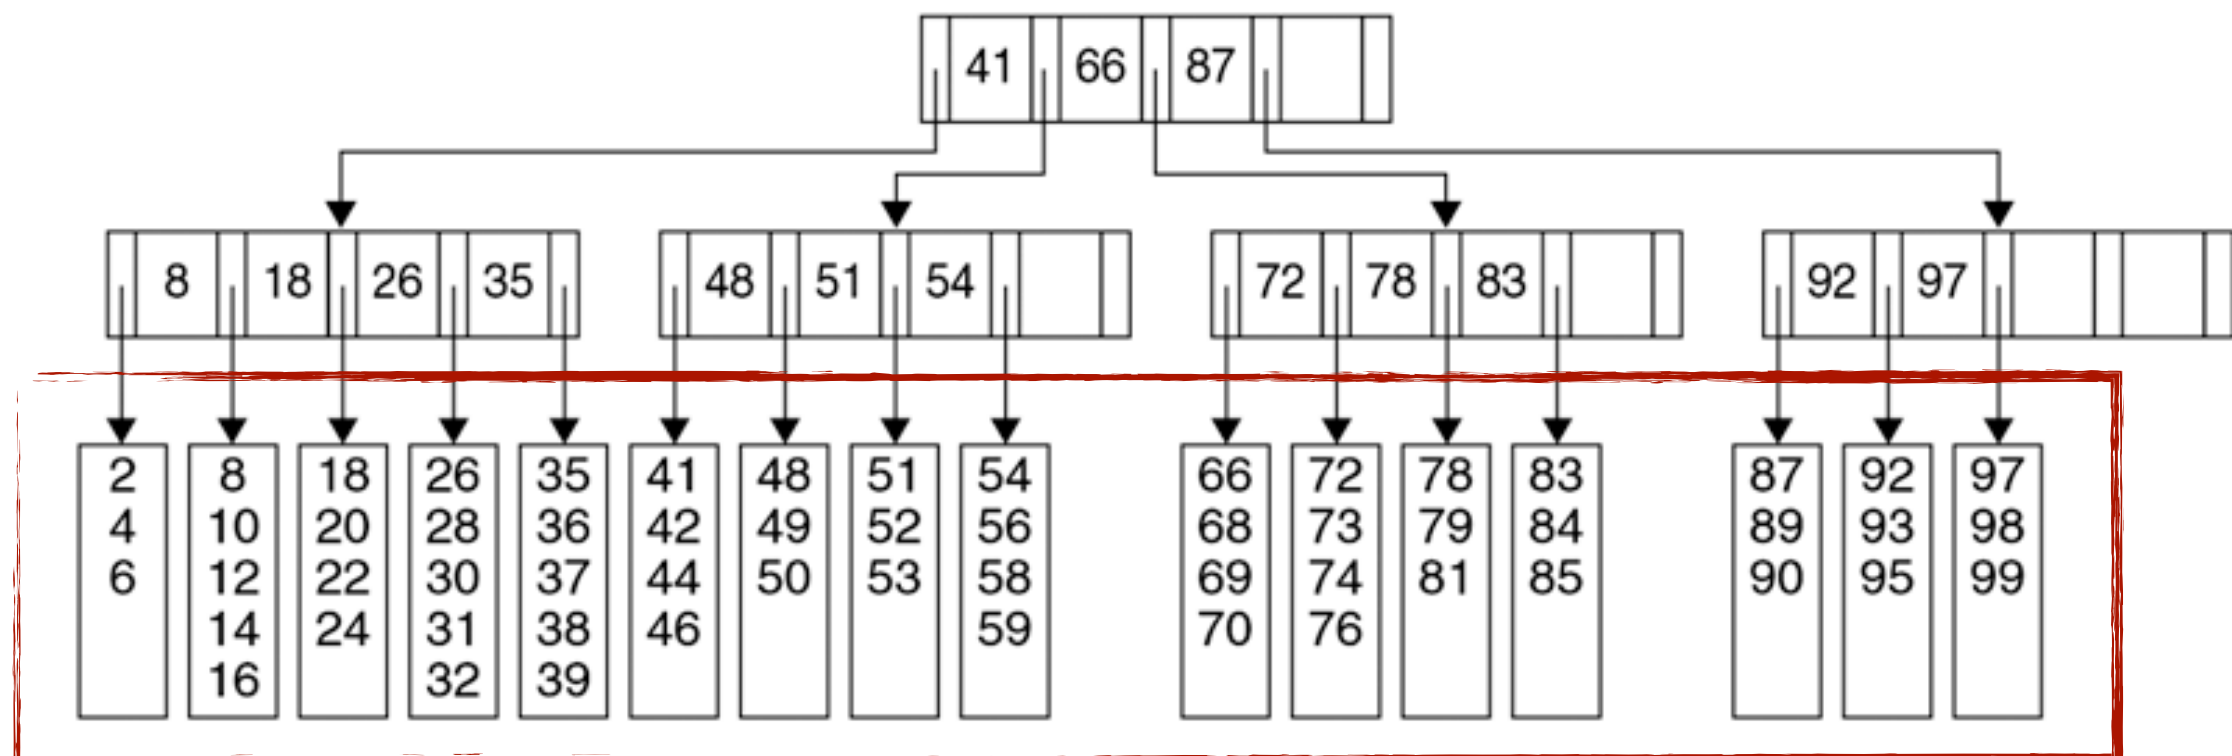
# B-Tree Example





# B-Tree Properties

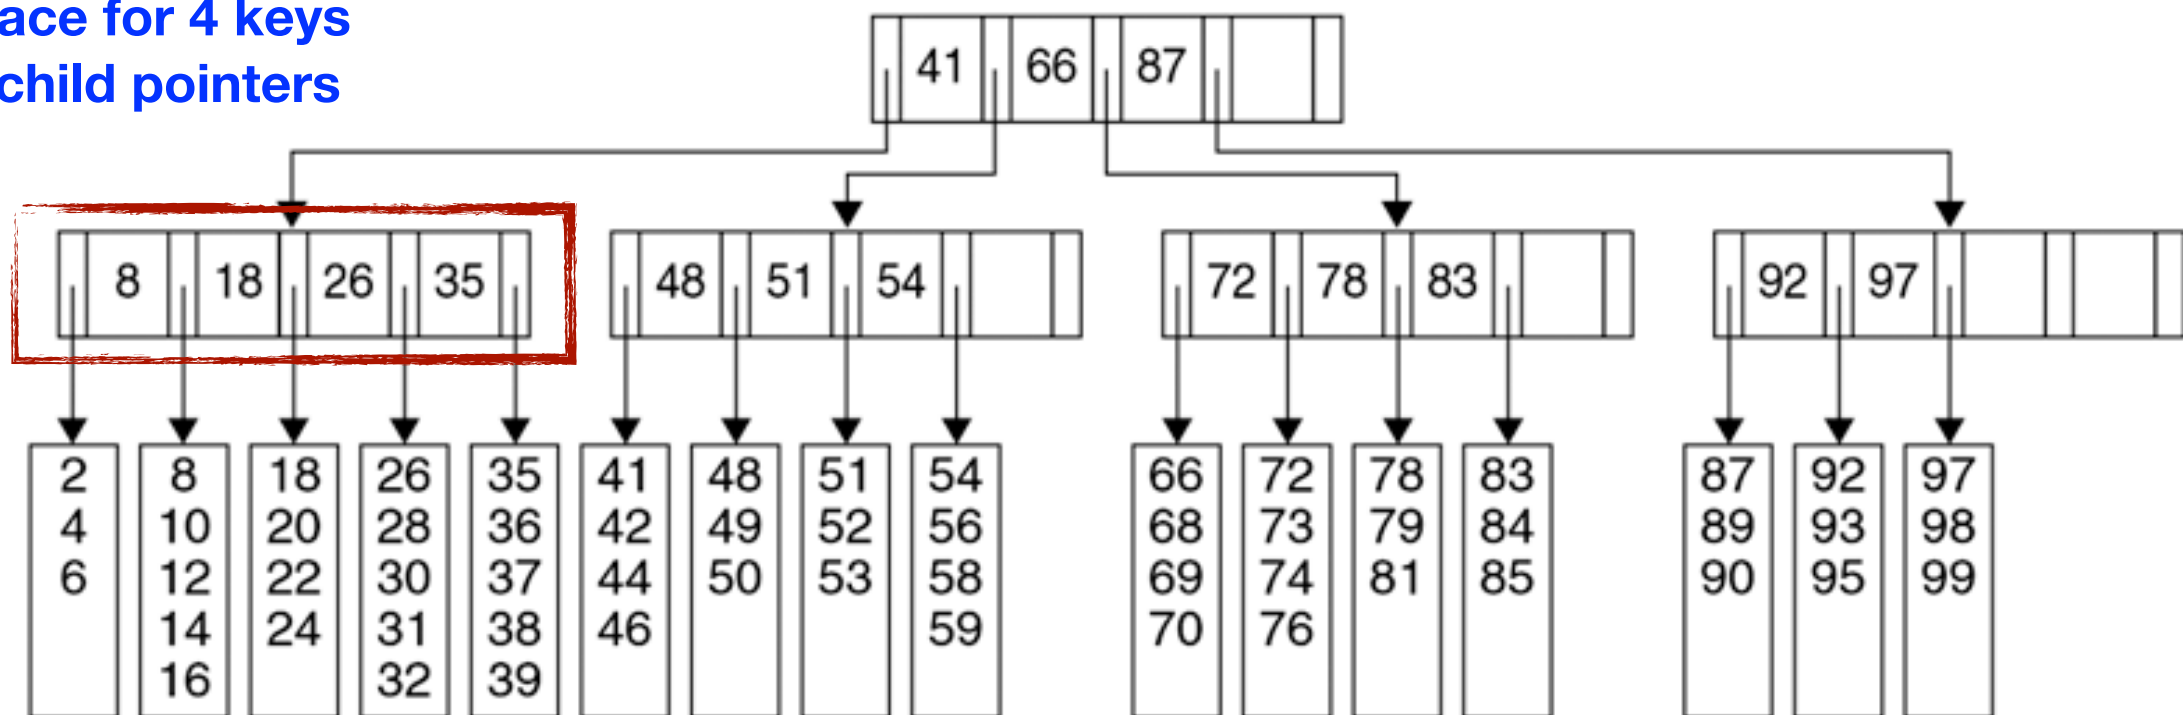
**(1) All data items are stored in the leaf nodes**



# B-Tree Properties

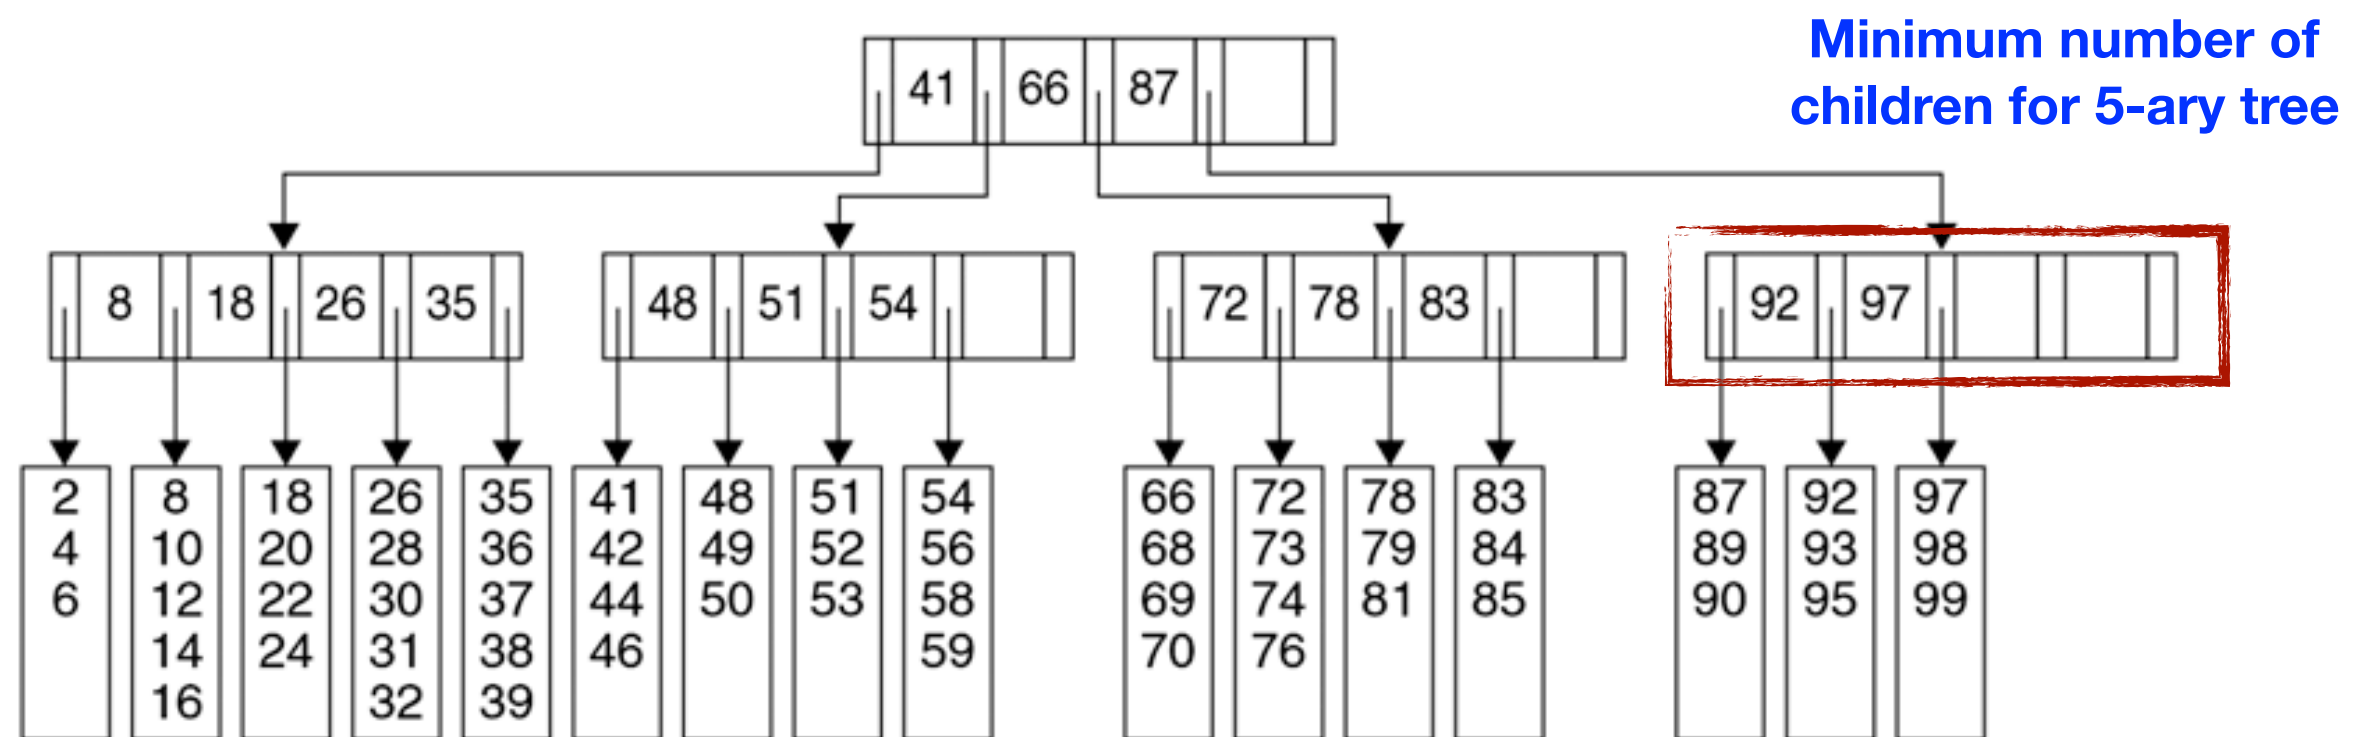
**(2) Non-leaf nodes store as many as  $M-1$  keys to guide the searching process; key  $i$  represents the smallest key in subtree  $i+1$**

5-ary tree, each node  
has space for 4 keys  
and 5 child pointers



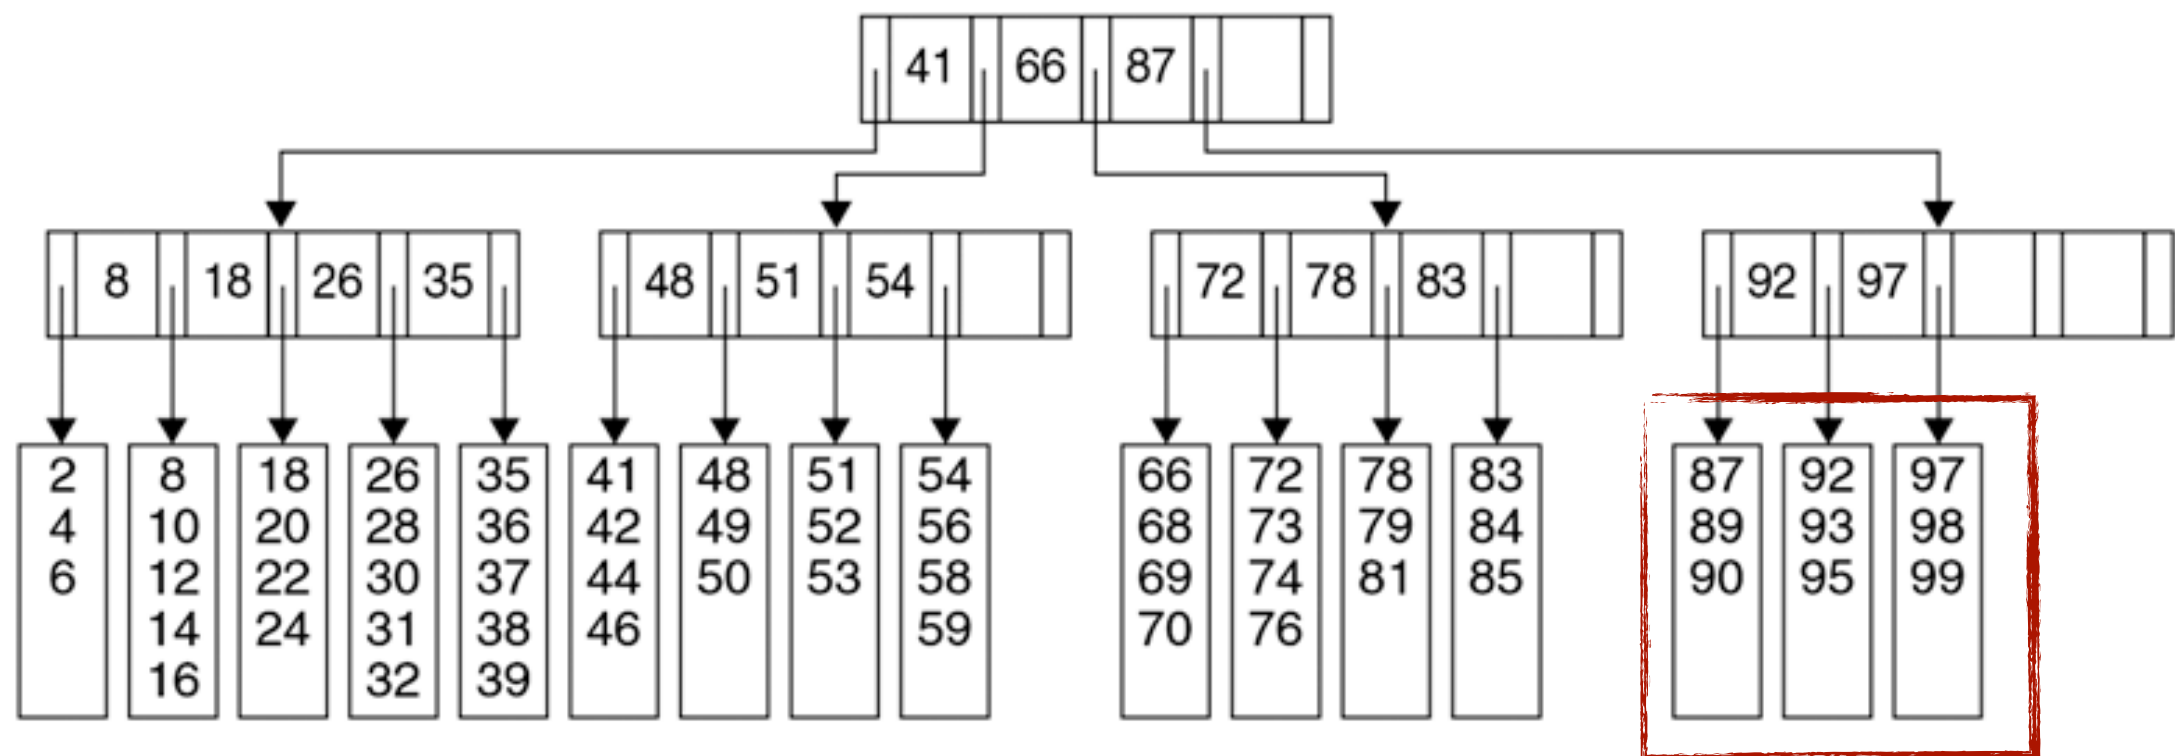
# B-Tree Properties

(4) All non-leaf nodes (except the root) have between  $\lceil M/2 \rceil$  and  $M$  children



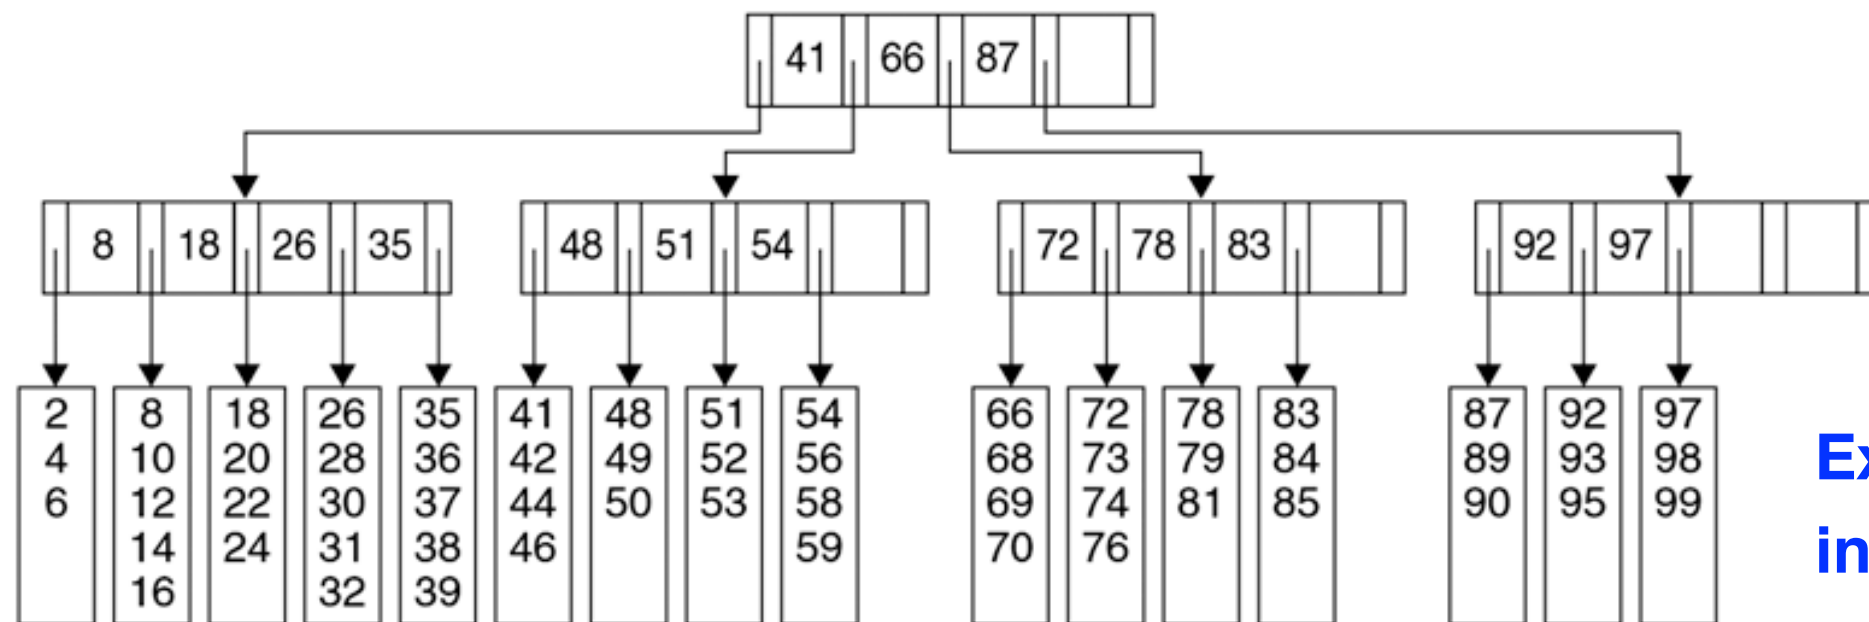
# B-Tree Properties

**(5) All leaf nodes are at the same depth and have between  $\lceil L/2 \rceil$  and  $L$  data items, for some value  $L$**

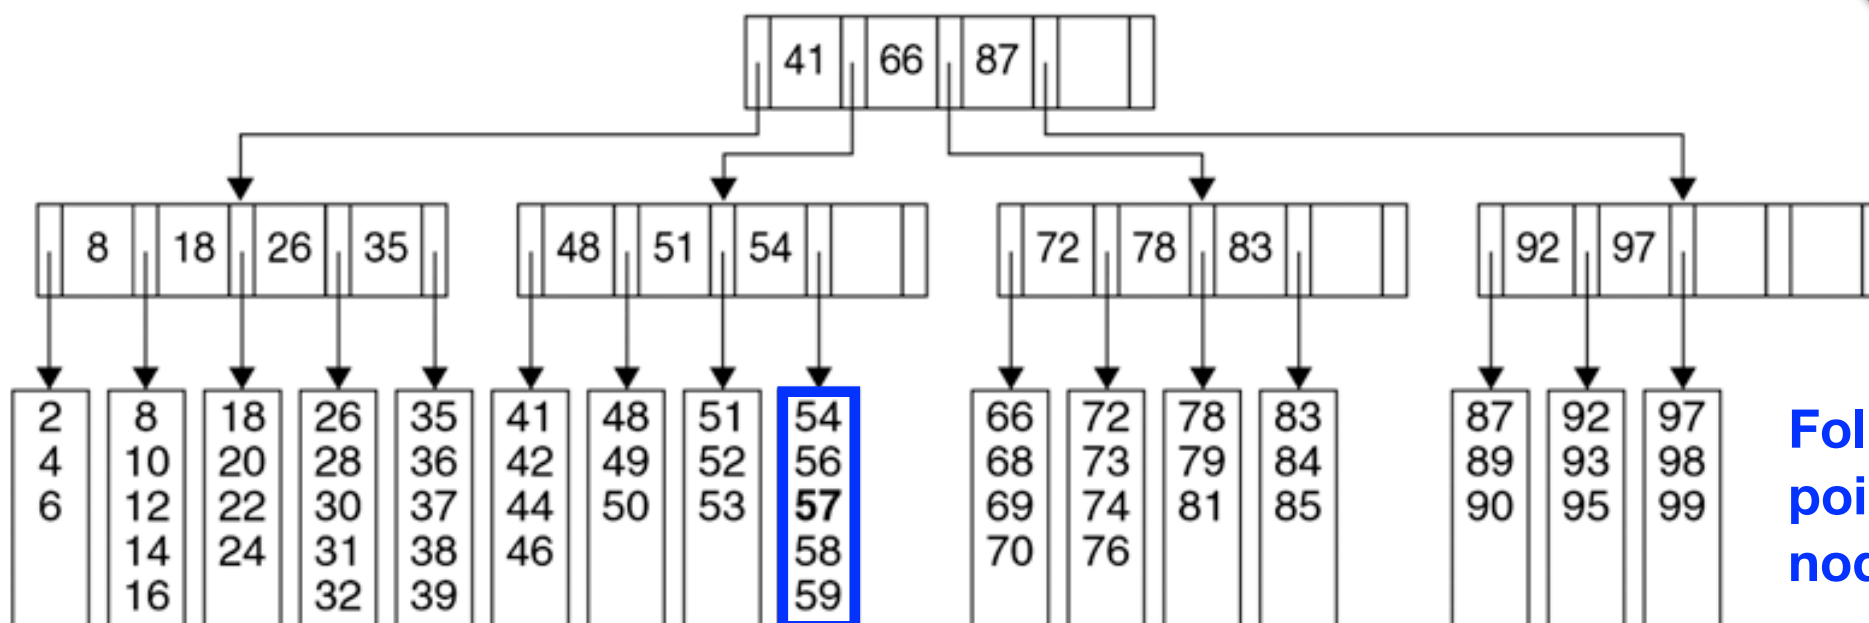


**In this example,  $L=5$  so each leaf node must have between 3 and 5 data items**

# B-Tree Insertion



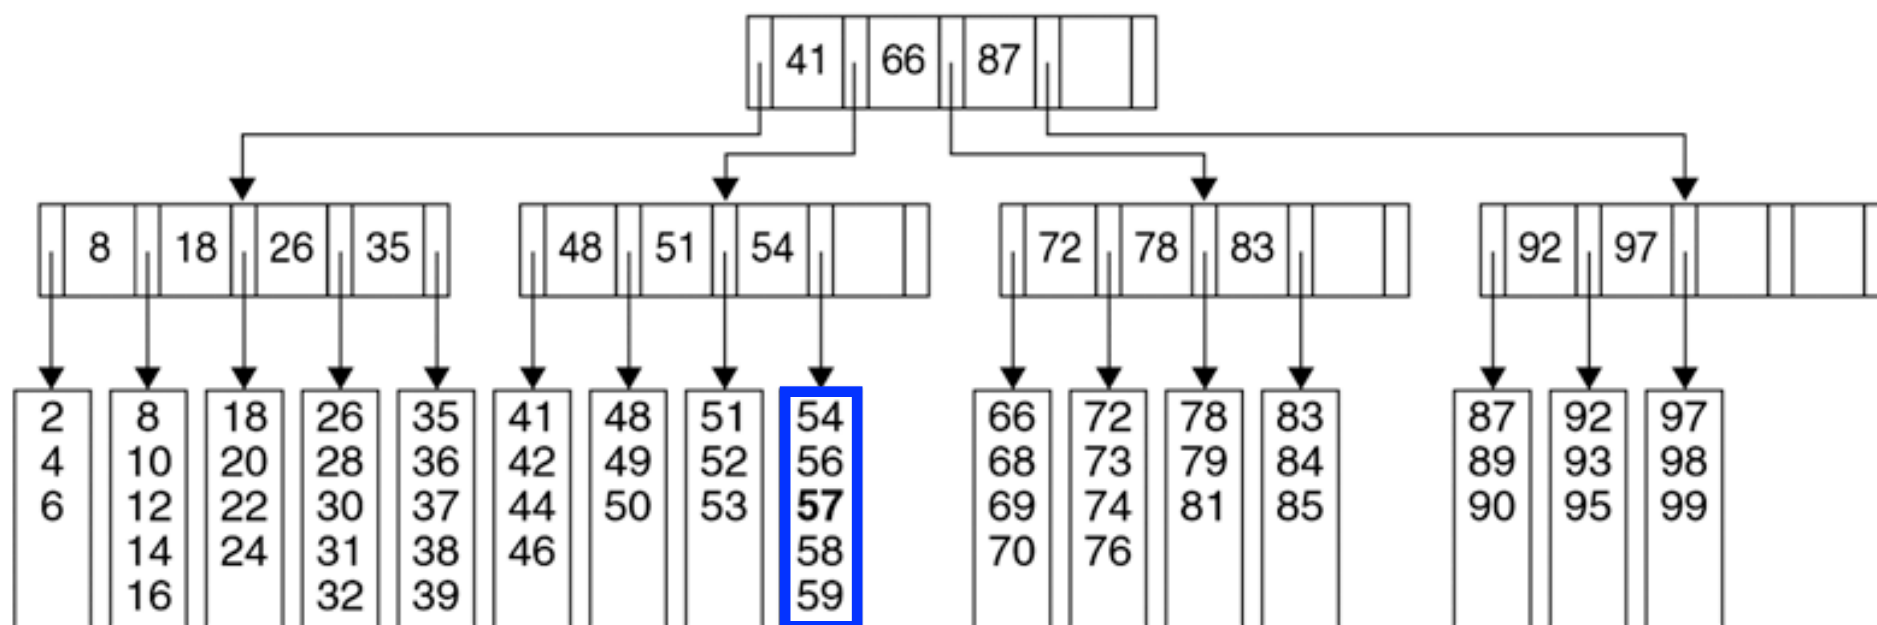
**Example of insertion:  
insert(57)**



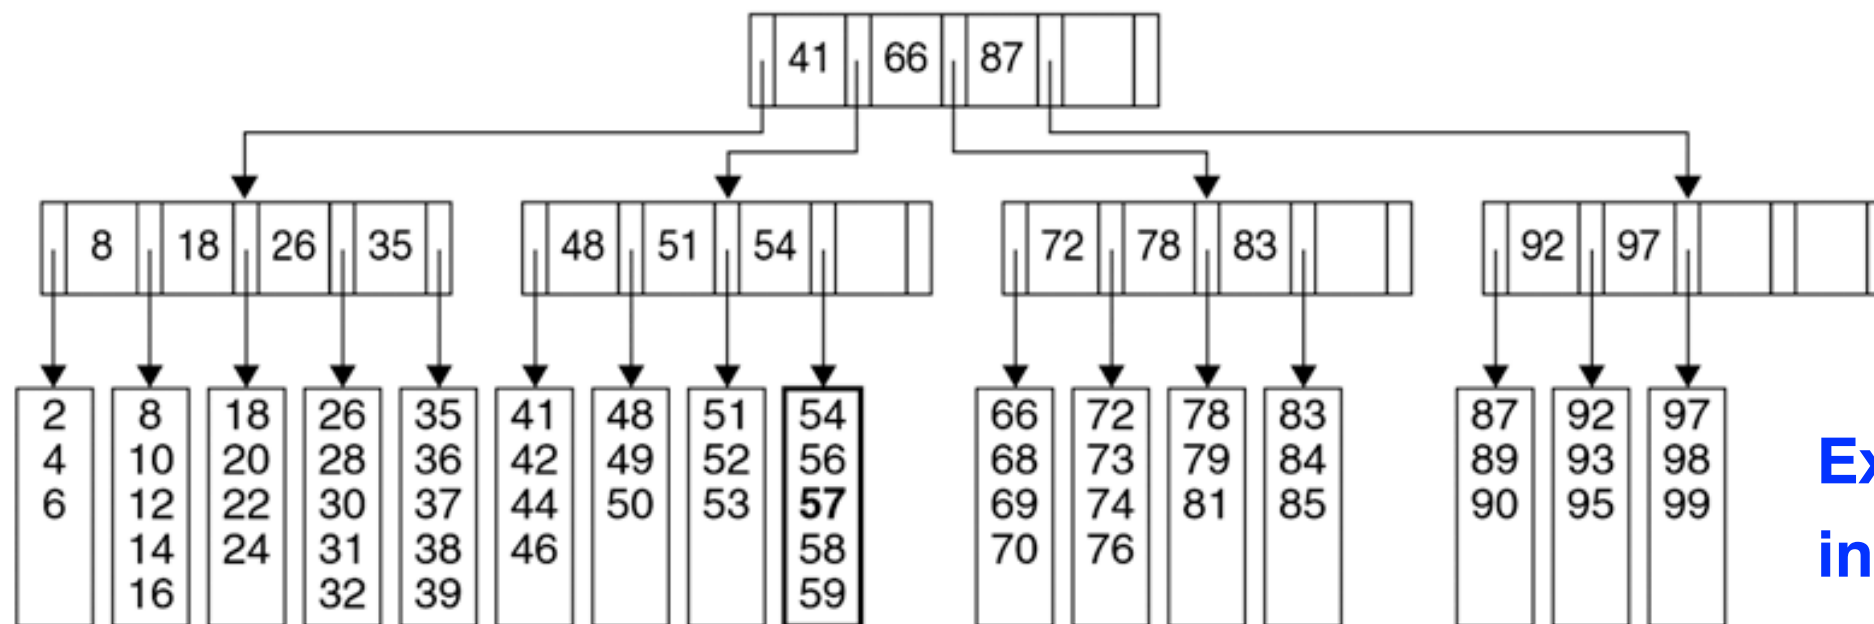
**Follow tree to insertion  
point and insert into leaf  
node if space is available**

# B-Tree Insertion

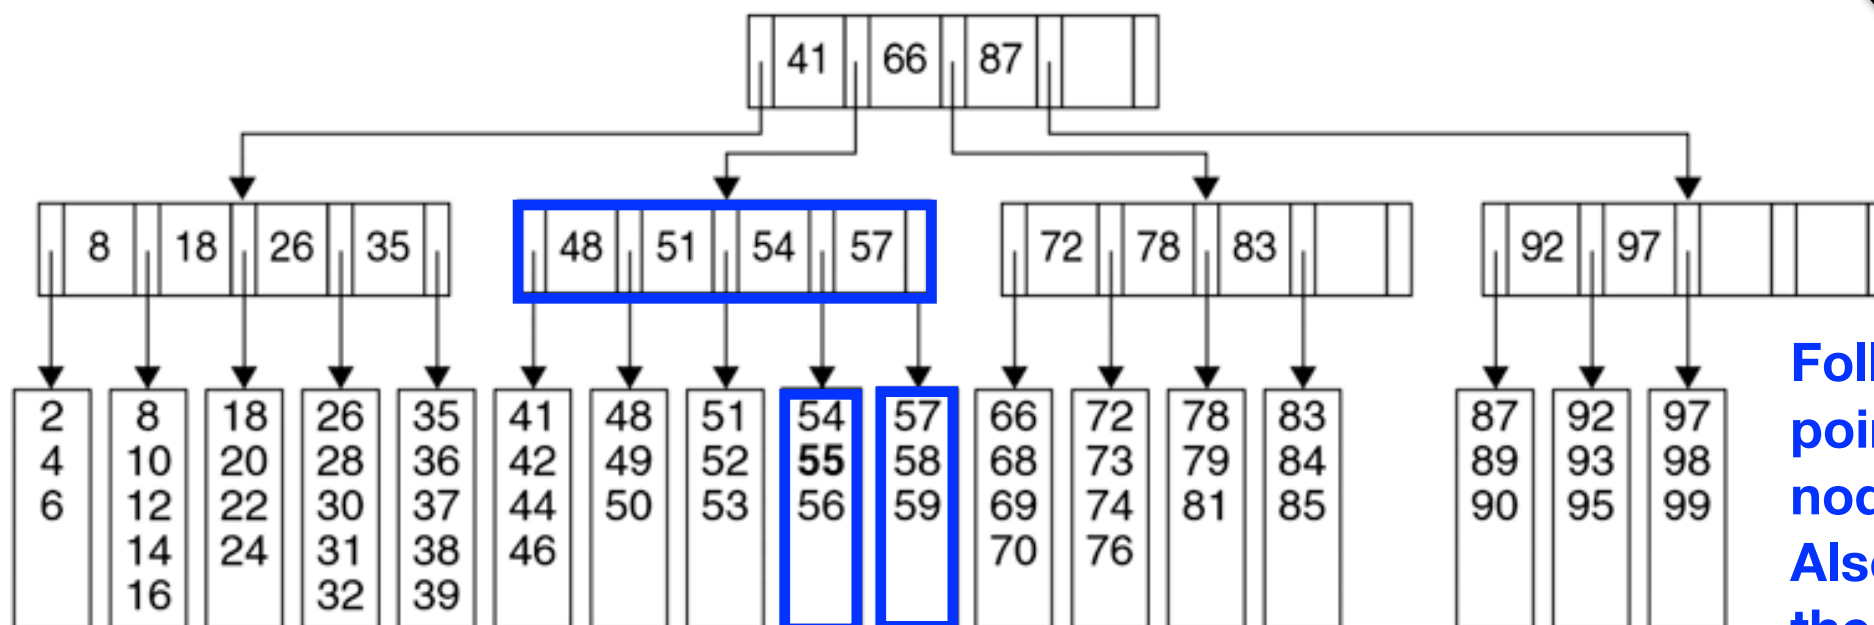
- **What if there is no space available in leaf node during insertion?**
  - The leaf node is split into two separate leaf nodes and the existing keys are redistributed between the new leaf nodes
- **Consider inserting the value 55 into the following tree**
  - The leaf node where 55 should be inserted is full



# B-Tree Insertion (With Full Leaf Node)



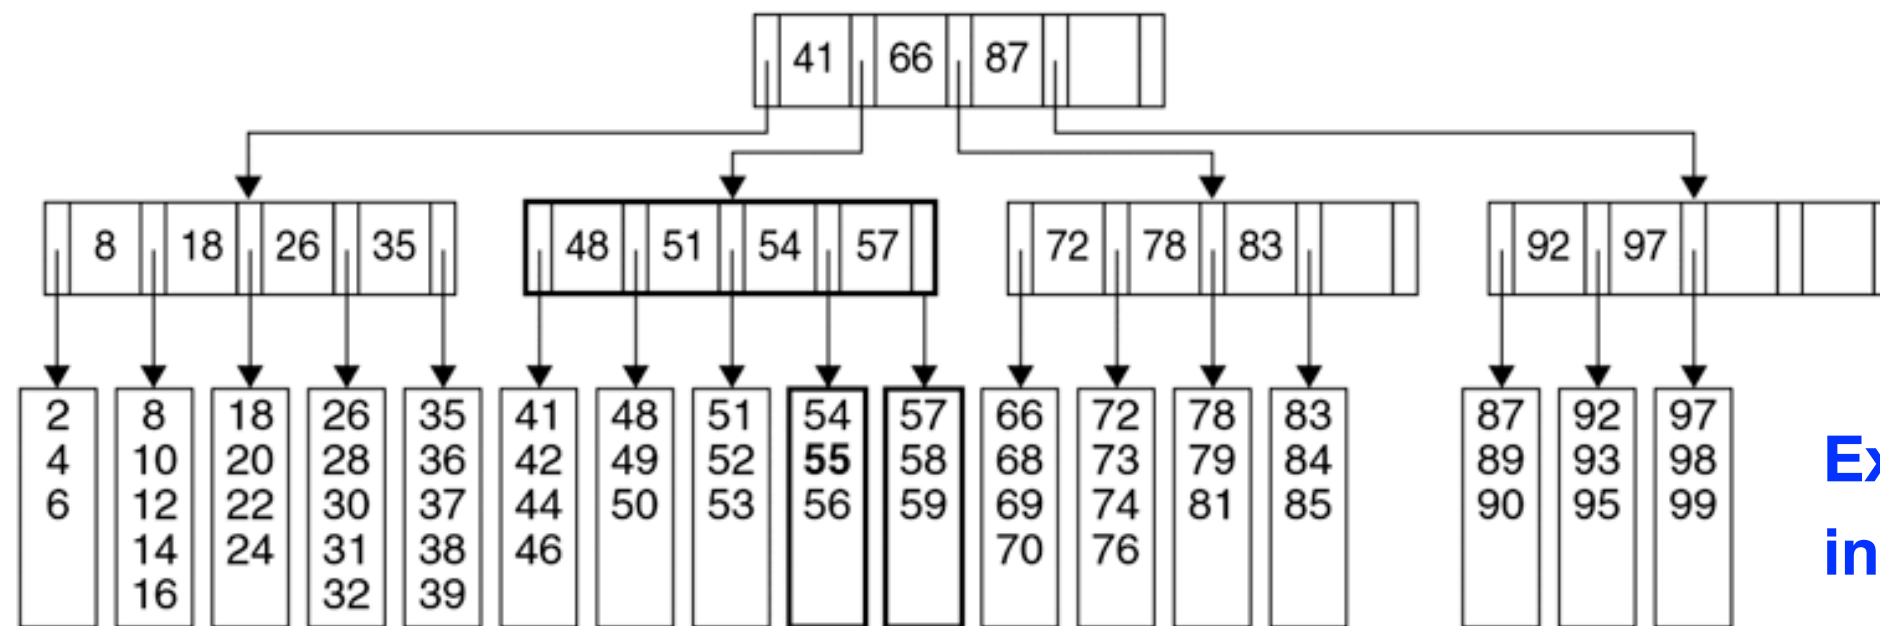
**Example of insertion:  
insert(55)**



**Follow tree to insertion  
point and split the full leaf  
node into two nodes.  
Also, update the keys in  
the parent node**

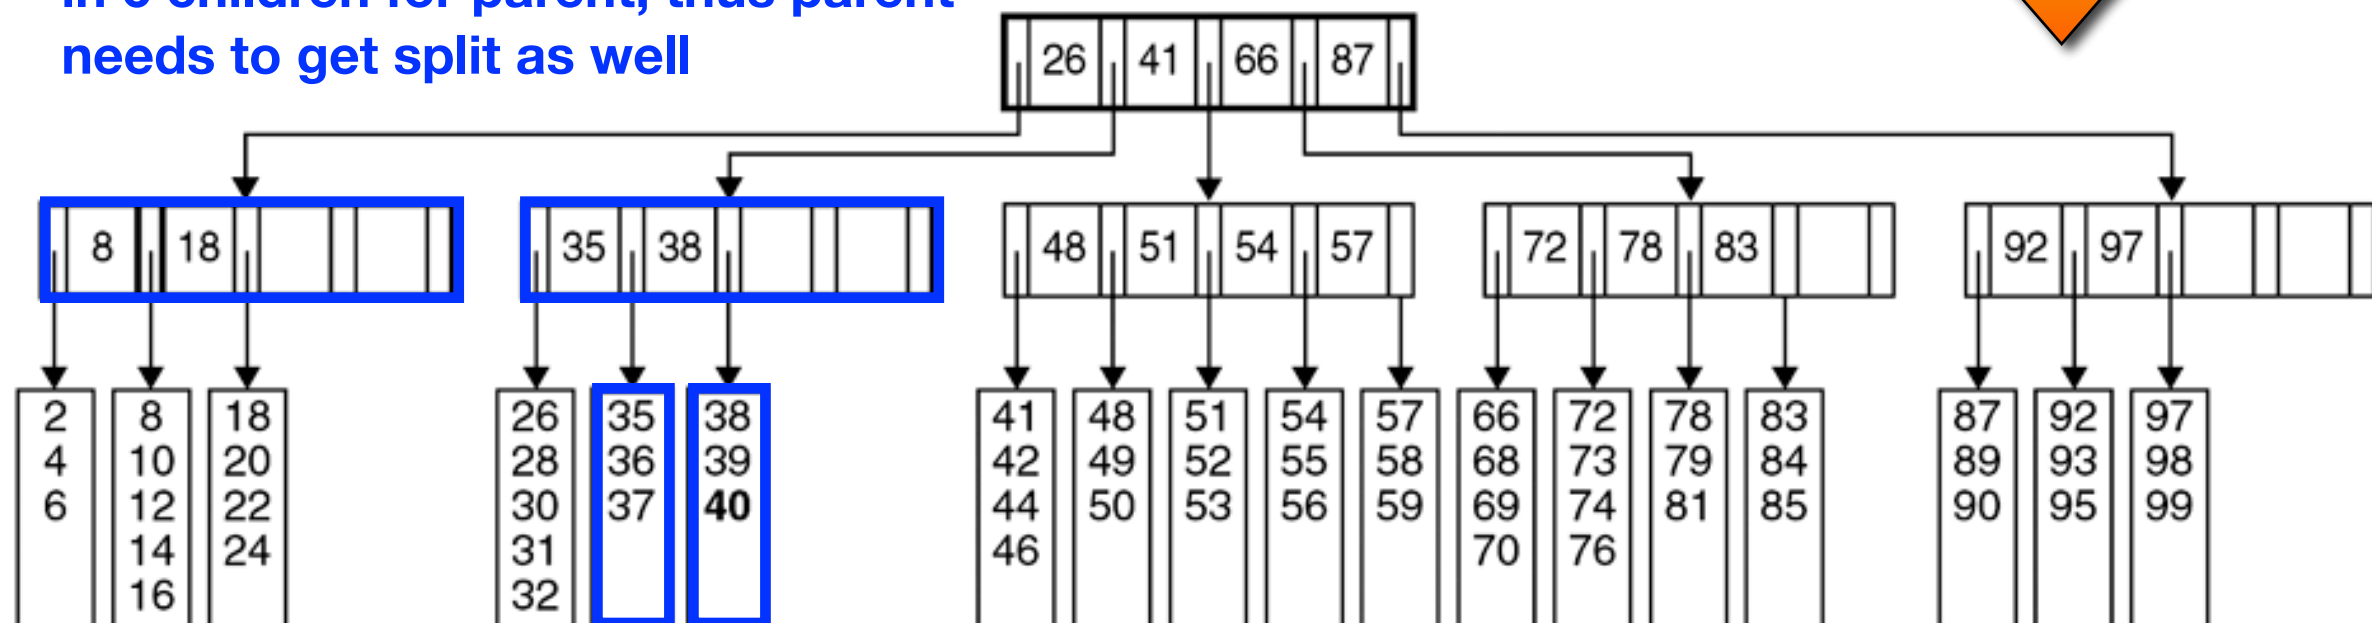


# B-Tree Insertion (With Full Leaf and Parent Nodes)



**Example of insertion:  
insert(40)**

Node will need to get split resulting  
in 6 children for parent, thus parent  
needs to get split as well





# B-Tree Insertion

---

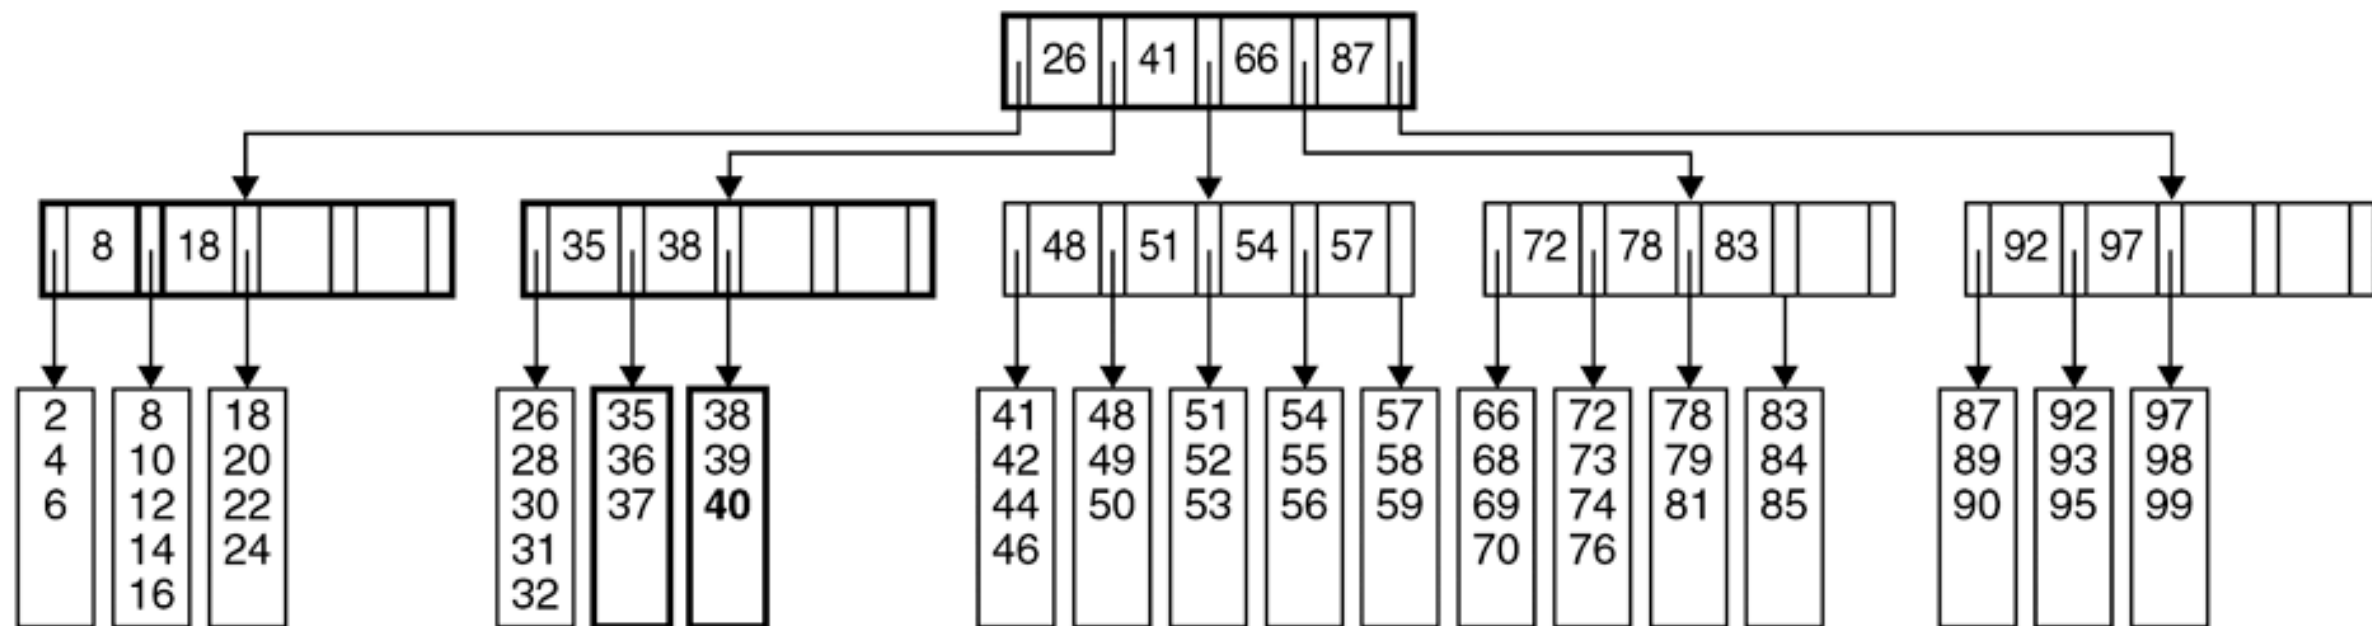
- **When a node is split, its parent gains a child**
- **If the parent has already reached its limit of children, then the parent must be split as well**
  - Must continue splitting all the way up the tree until a parent node is found that does not need splitting
  - If the root is split, then a new root must be created with the older, split root as its children (root is allowed to have only 2 children)

# B-Tree Deletion


---

- **To delete from a B-Tree, find the item that needs to be removed and remove it**
  - May cause violation of B-Tree rules if the leaf that the item is deleted from contains only the minimum number of items
    - Remember, each node must contain between  $\lceil L/2 \rceil$  and  $L$  data items
  - If a violation occurs, a node can gain the minimum number of items by adopting an item from a neighboring leaf node
    - If the neighbor is **ALSO** already at its minimum number of items, then the two leaf nodes can be combined to form a single full leaf node.
      - This may cause the parent to fall below the minimum number of children, in which case, the same adoption/combining strategy is used to satisfy the B-Tree properties

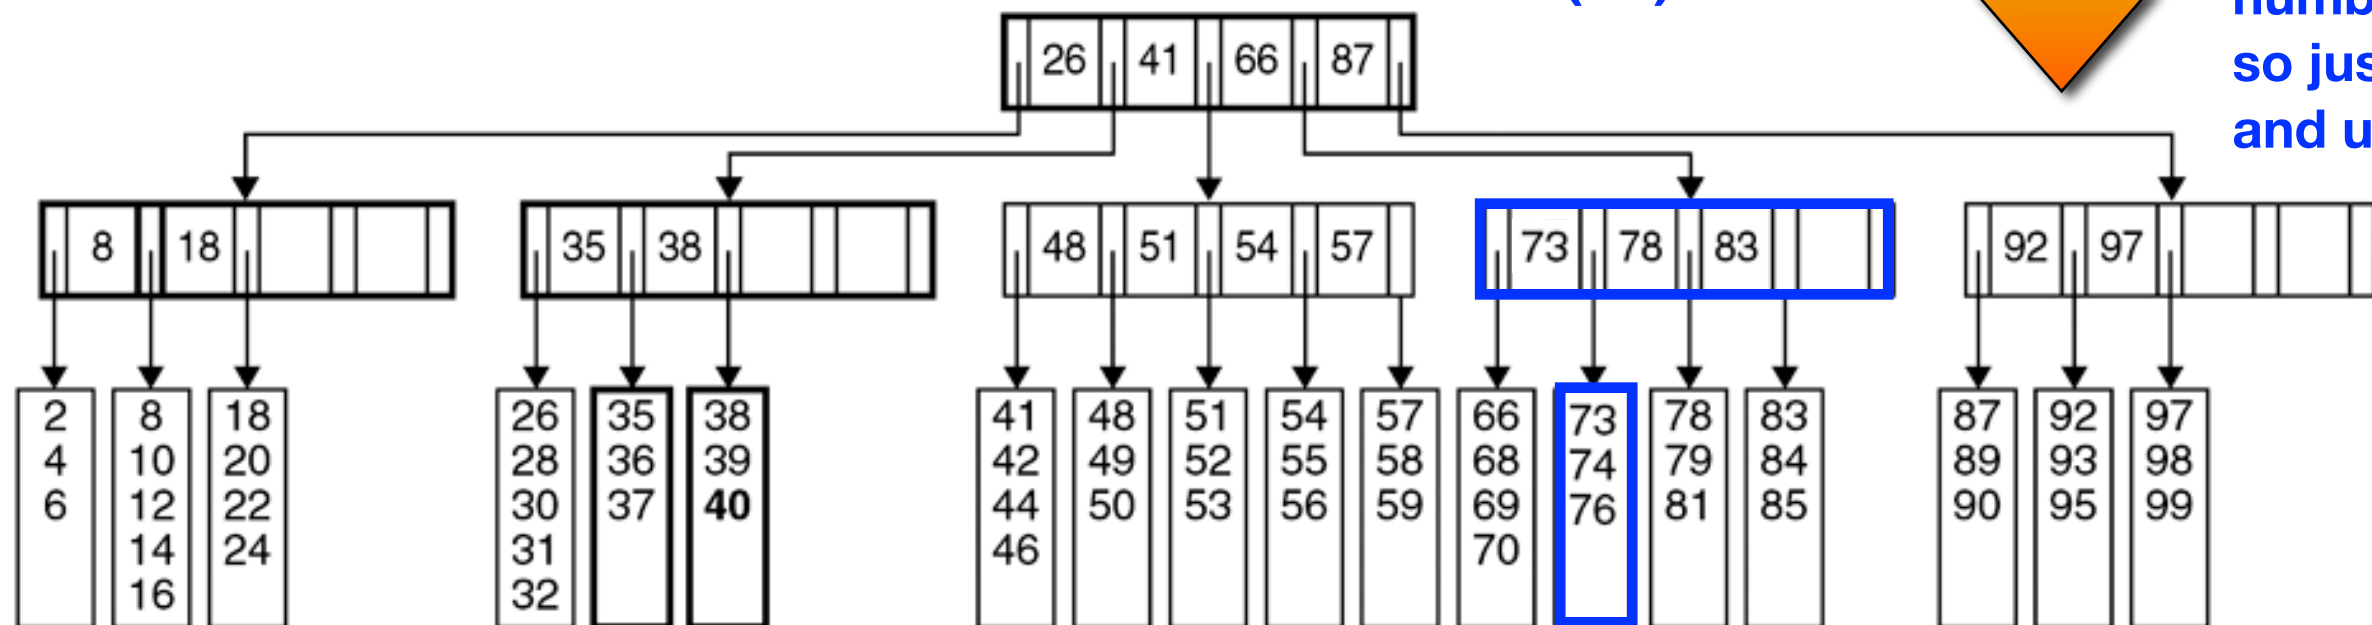
# B-Tree Deletion



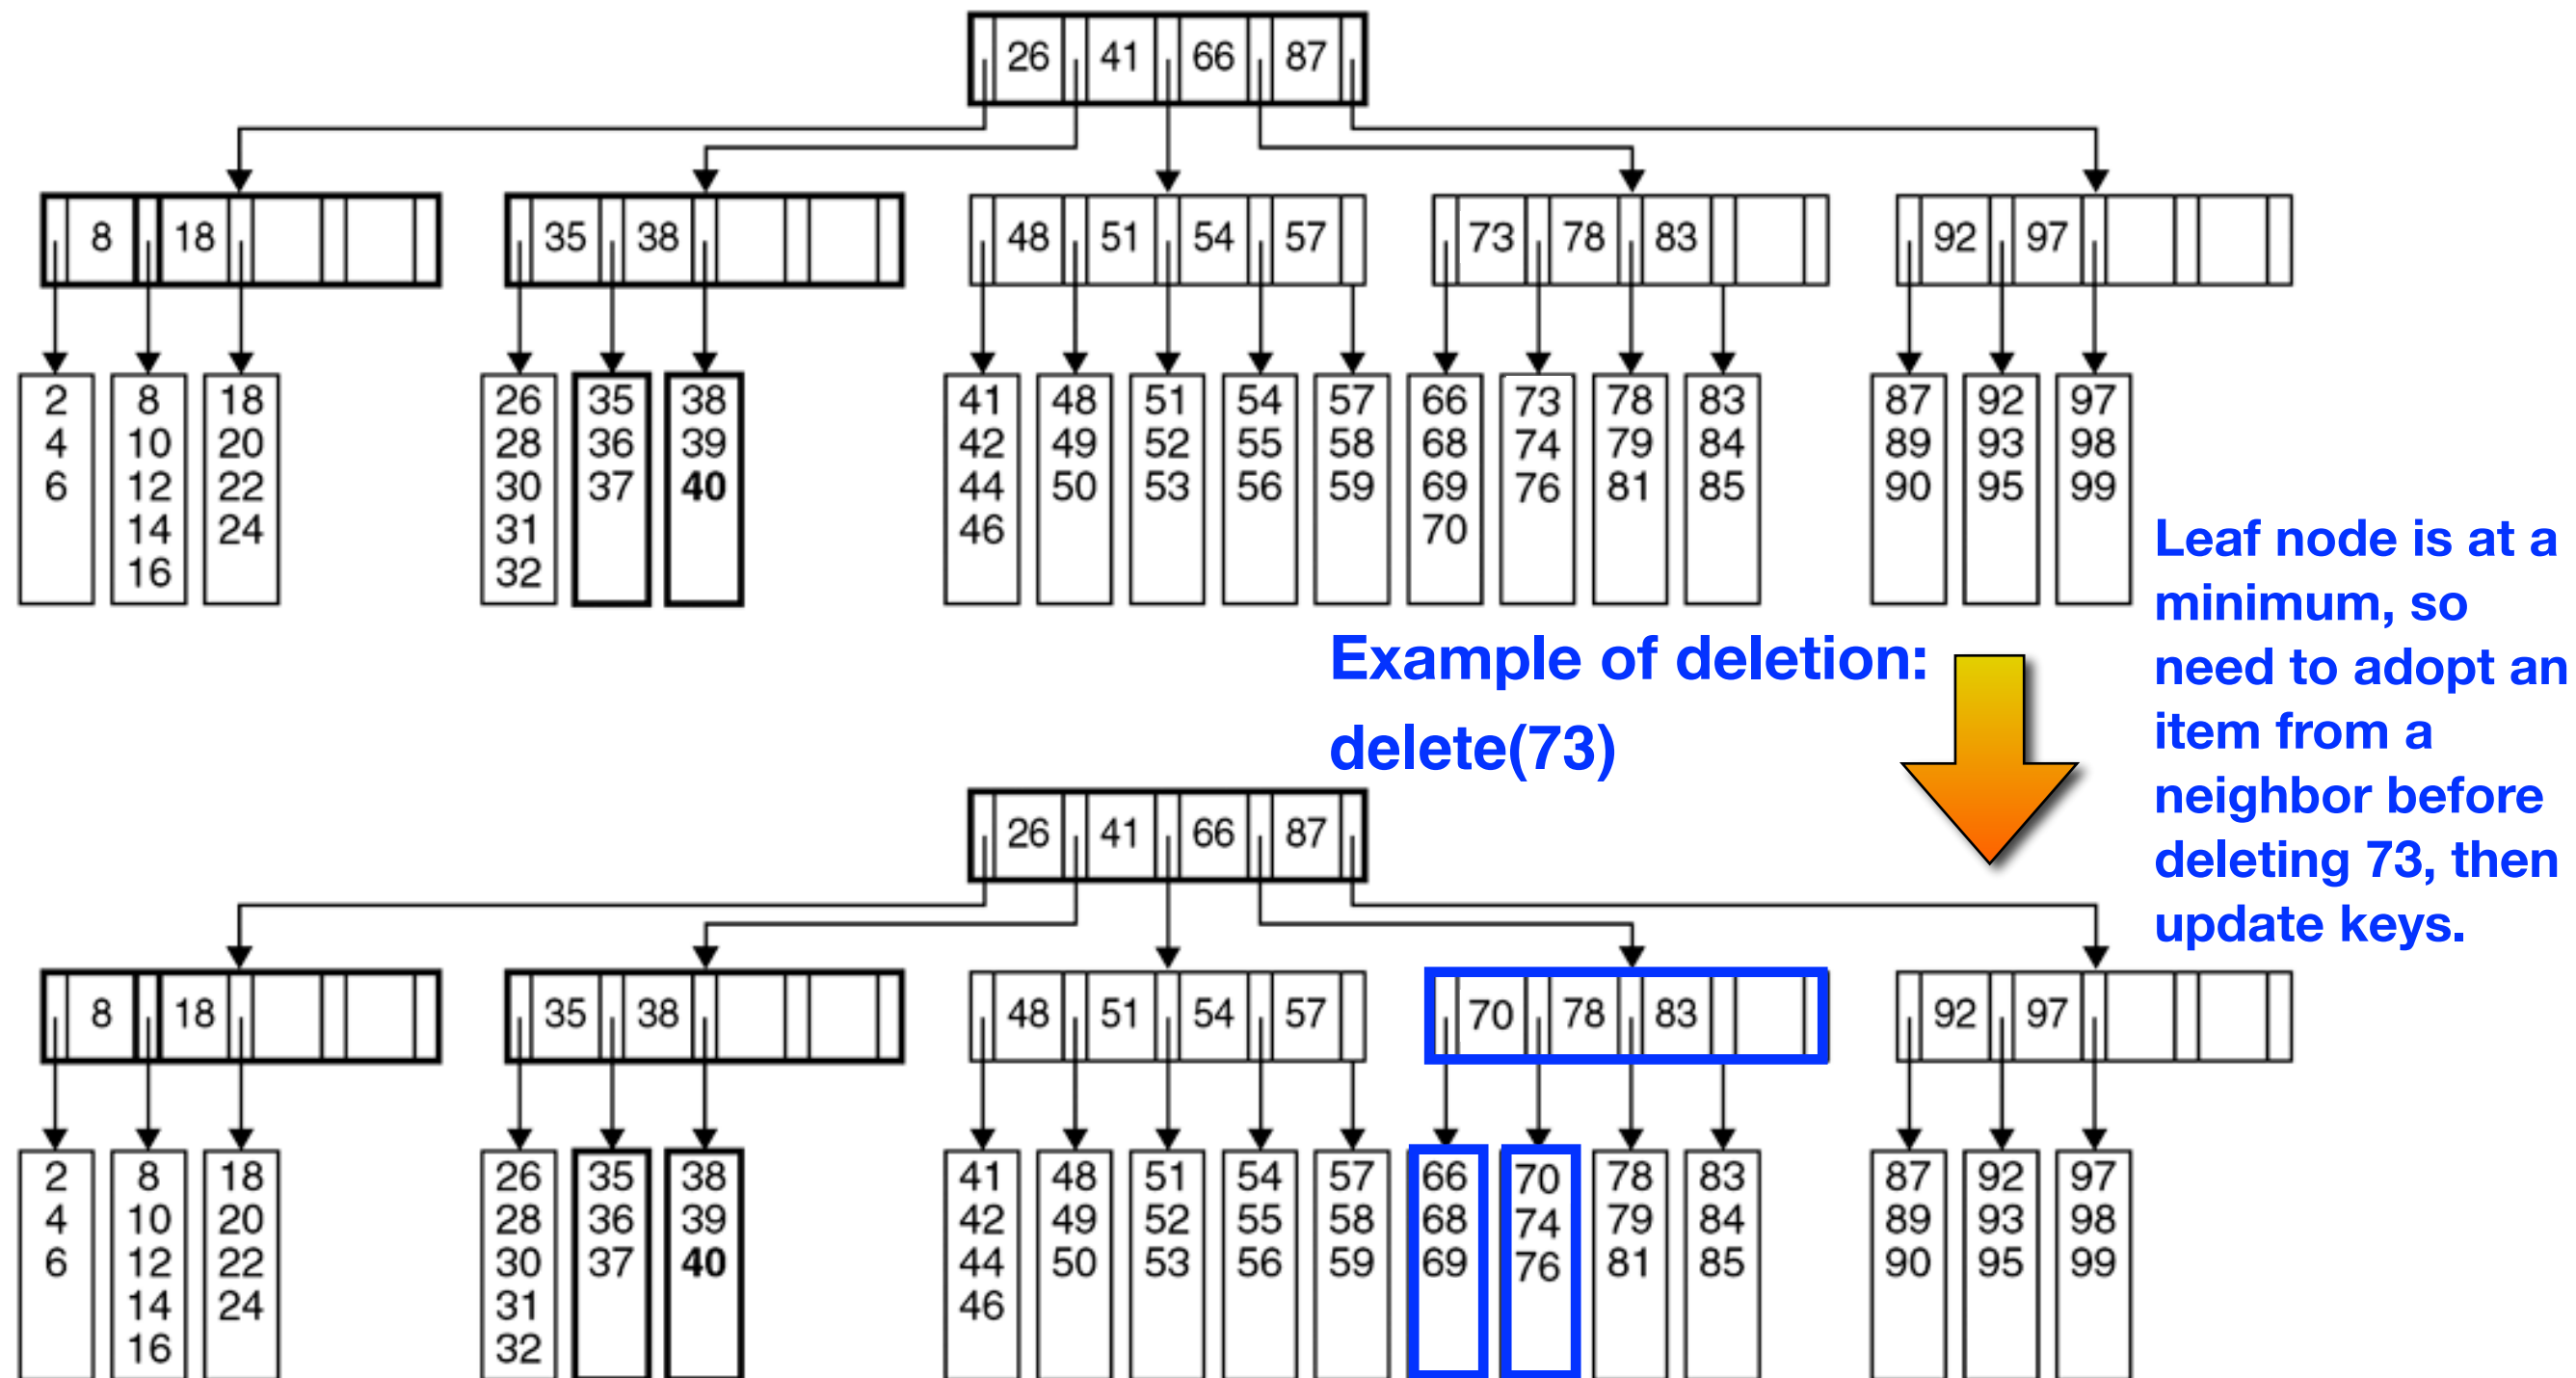
Example of deletion:  
**delete(72)**



Leaf node is not  
at a minimum  
number of items,  
so just delete 72  
and update keys.



# B-Tree Deletion



# B-Tree Deletion

